

Enhanced Sprint Plan with Detailed Developer Instructions

This revision deepens the guidance previously supplied by spelling out **day-by-day activities**, **acceptance criteria**, **deliverables**, and **hand-off checkpoints** for each of the four developers across all three sprints. All timings assume an 8-hour workday.

Sprint 1 – Foundation (Week 1)

Day	Dev 1 – Backend Lead	Dev 2 – Data Engineer	Dev 3 – Algorithm Dev	Dev 4 – Frontend / Integr
Mon (D1)	Create project repo, set up virtual-env, commit base pyproject.toml; stub qloo_client.py class; request Qloo API key and store in .env ^[1]	Fork repo, add data/grocery_catalog.csv, draft schema for associations.db; install pandas, sqlite3	Draft AssociationEngine interface; outline unit-test matrix in tests/test_association_engine.py	Scaffold Streamlit app (streamlit init shell with empty pages
Tue (D2)	Implement GET /recommendations/products wrapper with retry & exponential back-off; write happy-path & timeout unit tests	Build ETL script to import catalog into SQLite; publish Makefile target make_load_catalog	Write parsing helpers (parse_associations) incl. Pydantic models for response validation	Add sidebar brand, configure Tailwind CS for UI components
Wed (D3)	Finish POST /insights/preferences wrapper; benchmark latency (<2 s); push coverage badge	Develop batch job to call API for 1000 items/hr; use tqdm progress bar; persist to SQLite	Implement process_grocery_catalog() calling Dev 1 wrappers; unit tests ≥80%	Design "Association Network" placeholder event bus for real-time updates
Thu (D4)	Harden API module: circuit-breaker, 429 handling, ENV-switch for sandbox/production	Add incremental update logic using last_updated column; write DDL in database/schema.sql	Compute confidence scores; create top-5 association view; expose via associations.to_parquet()	Hook Parquet feed; render demo heatmap profile render time (<1 s)
Fri (D5)	Code review & merge Dev 1+3 PRs; publish package to internal PyPI	Profiling & optimization (indexing, WAL, PRAGMA tuning) – queries <100 ms for 10k rows ^[2]	Pair with Dev 4 to validate algorithm JSON contract	Implement layout_optimizer.py::generate_sl returning ≥200 placements ^[1] ^[2]
Sat (D6)	End-to-end smoke test: fetch 10 SKUs → DB → association JSON	Write nightly cron (schedule lib) for batch associations → push to branch feature/cron	Draft docstring examples; ensure mypy passes	Visual smoke test inside Streamlit; fix lint
Sun (D7)	Buffer / contingency; harden docs (docs/api.md)	Buffer / contingency	Buffer / contingency	Buffer / contingency

Sprint 1 Acceptance Criteria

- qloo_client.py returns valid JSON for sample SKUs in <2 s round-trip^[1].
- associations.db contains ≥10 000 product-pair rows, <100 ms avg read time^[2].
- layout_optimizer.generate_shelf_placements() outputs ≥200 placement dicts with schema {primary_product, companion_product, confidence, shelf_distance}.
- All modules achieve ≥80% unit-test coverage and pass on CI pipeline.

Sprint 2 – Core Features (Week 2)

Day	Dev 1 – Backend Lead	Dev 2 – Data Engineer	Dev 3 – Algorithm Dev	Dev 4 – Frontend / Integration
Mon (D8)	Branch feature/combo_generator; design dataclass Combo	Draft weekly KPI table design for reports; set up Jinja2 templates	Extend optimizer for section-level placement (optimize_store_sections)	Flesh out "Shelf Placements" page: DataTable + sortable confidence column
Tue (D9)	Implement high-confidence filter (≥0.8) and generate_weekly_combos(); unit tests	Build WeeklyReport model → YAML; accept CLI --week YYYY-MM-DD	Implement heuristic: crowding penalty + adjacency reward; parameterize via config.yml	Add Plotly Sankey diagram for product flow; anchor to section IDs
Wed (D10)	Integrate price API stub to suggest discount %; ensure configurable fallback	Aggregation SQL views for: new associations vs prior week; expose via Pandas	Refactor scoring into algorithms/scoring.py; publish ROC metrics notebook	Build "Combo Offers" UI with card layout; add CTA "Download JSON"

Day	Dev 1 – Backend Lead	Dev 2 – Data Engineer	Dev 3 – Algorithm Dev	Dev 4 – Frontend / Integration
Thu (D11)	Dockerize service <code>association_api</code> (FastAPI) to expose <code>/combos</code> endpoint	Schedule Airflow DAG (<code>weekly_report</code>) triggered Fridays 02:00; load to <code>reports/</code>	Benchmark optimizer runtime <180 s for 10 k SKUs; profile hotspots	Integrate Dag status feed into Streamlit (<code>st_autorefresh</code>)
Fri (D12)	Publish OpenAPI schema & Swagger docs; PR review	Unit tests for DAG idempotency; add Slack webhook on success/fail	Document tunable hyper-params; write Sphinx page	Optimize UI bundle size (<1 MB) by code-splitting; Lighthouse score ≥90
Sat (D13)	End-to-end test: <code>/combos</code> returns ≥20 offers with JSON schema	Generate sample HTML report; attach inline heatmaps via <code>base64</code>	Validate optimizer output with new offers; cross-check anomalies	Add responsive design fixes; QA on mobile viewport
Sun (D14)	Buffer / bug fixes	Buffer / bug fixes	Buffer / bug fixes	Buffer / bug fixes

Sprint 2 Acceptance Criteria

- 1. REST endpoint `/combos` delivers ≥20 weekly combos with ≥80% Qloo confidence^[1].
- 2. Airflow DAG runs automatically and writes HTML + assets in <5 min^[1] ^[2].
- 3. `optimize_store_sections()` optimizes 8 predefined sections with crowding penalty ≤5%.
- 4. Streamlit pages (“Associations”, “Placements”, “Combos”, “Reports”) load in <1 s (95th pct).

Sprint 3 – Interface & Polish (Week 3)

Day	Dev 1 – Backend Lead	Dev 2 – Data Engineer	Dev 3 – Algorithm Dev	Dev 4 – Frontend / Integration
Mon (D15)	Integrate modules via service mesh; write <code>docker-compose.yml</code>	Design OLTP schema for future POS ingestion; add migration scripts (Alembic)	Draft “association drift” detection notebook for future ML	Build dashboard homepage: KPI cards pulling from <code>/metrics</code>
Tue (D16)	Set up GitHub Actions: lint-test-build-deploy; add cache	Bulk load sample POS CSV; stress-test DB (10 k products, 1 M tx)	Author <code>docs/architecture.md</code> diagrams (Mermaid)	Implement auth layer (basic HTTP password) for dashboard
Wed (D17)	JMeter perf test: 50 RPS, target API p95 <2 s; tune gunicorn workers	Create read replicas & benchmarks; finalize indices	Polish docstrings; run <code>pydocstyle</code> ; generate API docs	Add interactive store map (Plotly) with placement overlays
Thu (D18)	Write 100+ integration tests; branch protection rules	Add DB backup script + cron; test restore in staging	Package CLI (<code>smo-cli</code>) for headless invocation; upload to PyPI test	UI polish: dark/light toggle, accessibility audit
Fri (D19)	Orchestrate deployment script <code>deployment/setup.sh</code> for fresh Ubuntu; 30-min install target ^[2]	Tune connection pooling; ensure query p95 <300 ms	Prepare CHANGELOG, semantic version tags (<code>v1.0.0-rc1</code>)	Cross-browser QA; fix edge cases in Safari
Sat (D20)	Final load test & harden security headers; OWASP ZAP scan 0 highs	Snapshot production DB; stage migration docs	Conduct dry-run hand-over demo; capture video	Add README badges (build, coverage, license)
Sun (D21)	Sprint review demo lead; ensure 99% uptime in staging since midnight	Present DB metrics; hand-off run-books	Present algorithm notebook & future roadmap	Live demo of full dashboard; gather feedback

Sprint 3 Acceptance Criteria

- 1. **Streamlit dashboard** offers 5 sections, authenticated, mobile-responsive, p95 load <1 s^[1].
- 2. System withstands 50 RPS for 30 min with uptime ≥99%, memory <1 GB/container^[2].
- 3. Fresh-host deployment completes in ≤30 min via `deployment/setup.sh`, no manual edits.
- 4. Full documentation (architecture, API, ops, user guide) published and linked from README.

Cross-Team Standards

Area	Standard
Branch naming	<code>feature/<slug></code> , <code>bugfix/<slug></code> , <code>hotfix/<slug></code>
Commit style	Conventional Commits (<code>feat:</code> , <code>fix:</code> , <code>docs:</code> ...)

Area	Standard
Pull Requests	1 reviewer min; include screenshot/benchmark where relevant
Testing	pytest + coverage; failing tests block merge
CI	Lint → Unit tests → Build image → Push to registry → Deploy to staging
Definition of Done	Code, tests, docs, review, performance budget, acceptance tests pass

Sprint Ceremonies & Artefacts

Ceremony	Time-box	Artefact
Sprint Planning	2 h Monday 09:00	JIRA sprint board with story points and owner
Daily Stand-up	15 m 09:00	Slack summary thread ("Yesterday / Today / Blockers")
Mid-sprint Demo	30 m Wednesday	Loom video or live Zoom showing progress
Sprint Review	1 h Sunday	Working software demo, KPI dashboard snapshot
Retrospective	1 h Sunday	"Start / Stop / Continue" Miro board; action items assigned

By following these **granular, outcome-oriented instructions**, each developer knows exactly **what to code each day, how to validate it, and which measurable criteria define success**—ensuring visible, high-quality deliverables at the close of every sprint.



1. Qloo-Powered-Supermarket-Layout-Optimizer-Implementation-Plan.pdf
2. Supermarket-Product-Association-Layout-Optimization-Requirements-Document.pdf